

Benchmark Collection Assistant

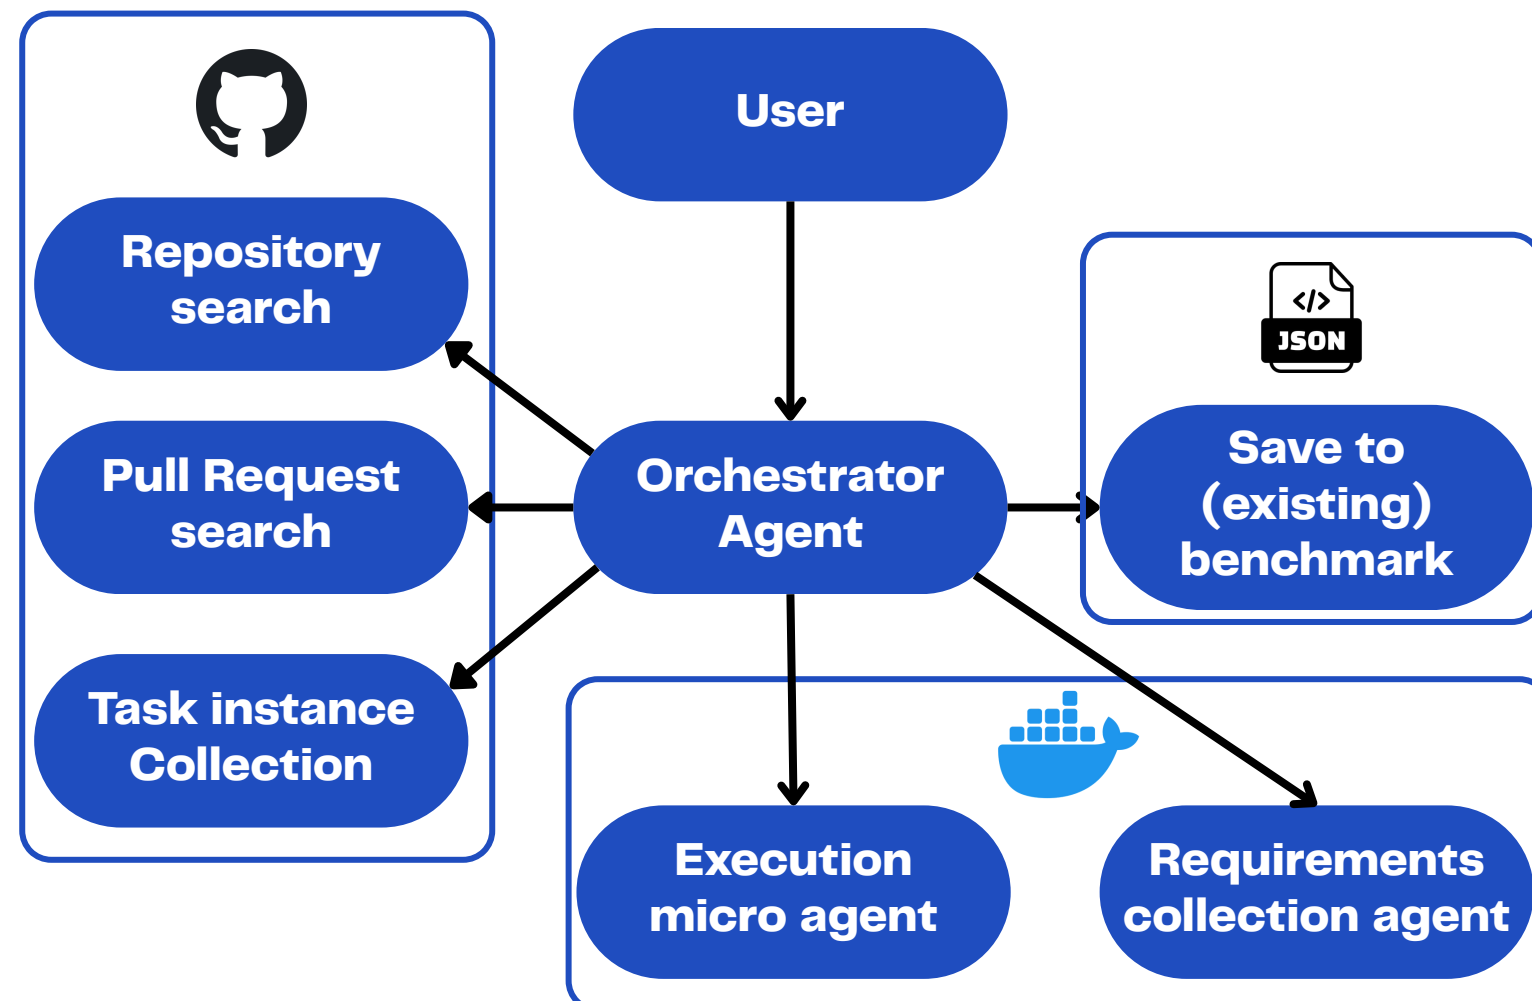
Problem:

Collecting benchmarks for evaluating AI features is time-consuming and (due to data leakage) a never-ending task

Solution:

Automate benchmark collection with an AI Agent!

Agent Architecture & Tools



Features

- 🔍 Customizable **Repository** Search, filter by:
 - Language
 - Build system
 - Specific Maven **m** & Gradle 🐘 dependencies
 - Many more...
- 🔍 Customizable **Pull Request** Search, filter by:
 - SWE-bench compatibility
 - Complexity of changes
- 🐳 Project setup in **Docker** container
 - Java ☕ & Kotlin 🇰🇷 support
- 🐱 **SWE-bench** format benchmark creation & extension

Results so far

🇰🇷 Kotlin benchmark integrated into the @AI Assistant Evaluation pipeline
☕ **Automatically** found Java version for >80 tasks in existing @AI Assistant benchmark
🐱 Generally observed a **dramatic speedup**

Contributing

Our agent is well-documented and extendable, so you can add your own features as well!

Upcoming Features

- More specific Repository filters
- Dependency search for more build systems
- More specific Pull Request filters
- Project setup support for **more languages**
- Saving **dockerfiles** upon project setup
- Parallel project setup
- **Metamorphic transformers** for robustness evaluations
- Semantic similarity-based **task clustering**
- **Your suggestions!**

Knowledge Base



Source Code



Processes

